

## EE450 Socket Programming Project, Fall 2013

**Due Date: Sunday November 24th, 2013 11:59 AM (Noon)**

**(The deadline is the same for all on-campus and DEN off-campus students)**

**Hard deadline (Strictly enforced)**

The objective of this assignment is to familiarize you with UNIX socket programming. This assignment is worth **10%** of your overall grade in this course.

**It is an individual assignment and no collaborations are allowed. Any cheating will result in an automatic F in the course (not just in the assignment).**

If you have any doubts/questions please feel free to contact the TAs and cc professor Zahid, as usual. Before that, make sure you have **read the whole project** description carefully.

### **Problem Statement**

In this project you will be simulating a resource-discovery / file-sharing network using a hybrid architecture with TCP and UDP sockets. The project has three major phases: 1) registration, 2) resource discovery, and 3) file-sharing. In phases 1 and 2, all communications are through UDP sockets. In phase 3 however, all the communications are over TCP connections i.e. through TCP sockets. The main components of this network architecture are: 1) three file servers, who have the files to be served to the clients, 2) one directory server, who operates as a coordinator between the clients and the file servers, and 3) two clients who contact the directory server to ask where (i.e. in which file server) to find the file they want to get.

### **Input Files Used**

The files specified below will be used as inputs in your programs in order to configure the state of the network.

1. **topology.txt**: This input file contains information about the cost of accessing each file server from each client (i.e. the network topology). Since there are 2 clients and 3 file servers, **topology.txt** has only 2 lines and each line contains the costs (i.e. a positive integer) to reach each file server from a client, separated by a single space. For instance, an example of the contents of the topology file can be:

2 45 87

11 3 27

This means that client 1 can reach file server 1 with a cost of 2, file server 2 with a cost of 45, and file server 3 with a cost of 87. Similarly client 2 can reach file server 1 with a cost of 11, file server 2 with a cost of 3, and file server 3 with a cost of 27.

2. **resource.txt**: This input file contains information about the resources stored in each file server. There are a total of 2 resources (i.e. files) available, called **doc1** and **doc2**. For instance the contents of **resource.txt** may be the following:

```
File_Server1 2 doc1 doc2
```

```
File_Server2 1 doc2
```

```
File_Server3 2 doc1 doc2
```

In each line above, you can find the name of the file server followed by its number of resources, as well as their names. You are given a sample **topology.txt** and **resource.txt** file to test your code. However, when your project is graded, the TA may use different input files (i.e. different costs, and different distribution of the two files among the file servers) to test your project. However, the format of the files, and the filenames would be the same.

### **Source Code Files**

Your implementation should include the source code files described below, for each component of the network.

1. **Directory Server**: You must use one of these names for this piece of code: **directory\_server.c** or **directory\_server.cc** or **directory\_server.cpp** (all small letters). Also you must call the corresponding header file (if any) **directory\_server.h** (all small letters). You must follow this naming convention. This piece of code basically represents the directory server in the project.
2. **File Server**: The name of this piece of code must be **file\_server.c** or **file\_server.cc** or **file\_server.cpp** (all small letters) and the header file (if any) must be called **file\_server.h** (all small letters). You must follow this naming convention. In order to create three file servers in your network, you can use the `fork()` function inside your `file_server` code to create child processes. However, if you are not familiar with `fork()`, you must create 3 instances of this code namely, **file\_server1.c**, **file\_server2.c** and **file\_server3.c** or **file\_server1.cc**, **file\_server2.cc** and **file\_server3.cc** or **file\_server1.cpp**, **file\_server2.cpp** and **file\_server3.cpp**.
3. **Client**: The name of this piece of code must be **client.c** or **client.cc** or **client.cpp** (all small letters) and the header file (if any) must be called `client.h` (all small letters). You must follow this naming convention. In case you are not using `fork()`, you must create 2 instances of this code namely, **client1.c**, **client2.c** or **client1.cc**, **client2.cc** or **client1.cpp**, **client2.cpp**.

### **Phase 1: (Registration)**

In the first phase of the project, you will implement the registration process of the file servers to the directory server. Specifically, the file servers connect to the directory server through UDP and inform the directory server of their associated TCP port numbers which the clients would be using in phase 3 to connect to them and obtain the required resources (i.e. files of interest). The directory server is responsible for creating a new text document called **directory.txt**, which contains the registration information (more details to be given below).

### **Phase 2: (Resource Discovery)**

In the second phase, the clients place a request for the file of interest to the directory server. The directory server will check the **resource.txt** file to find out which file server has the requested file. In case of multiple file servers having the requested file, the directory server looks into the **topology.txt** file to find the nearest file server (i.e. the one with the lower cost) to the client. The directory server will now send back the file server name, as well as its associated TCP port number to the client. In this phase, all communications take place using UDP.

### **Phase 3: (File Sharing)**

In this phase, clients establish a TCP connection with the file server by using the file server's TCP port number that was obtained from the directory server in phase 2. After the file server receives the request for the file from the client, it will send the requested file to the client.

### **More Detailed Explanation**

#### **Phase1:**

In this phase, the directory server creates a UDP socket in order to get the registration information from the file servers (see Table 1 for the static UDP port number to be assigned to the directory server, and the file servers). File servers will now communicate with the directory server using the directory server's UDP static port number. Each file server sends the following message to the directory server:

`File_Server# TCP_Port_No`

As an example, the message could be:

`File_Server1 22080`

Directory server must also create a text file called **directory.txt** (you must use this name for the

text file). Please note that the directory server creates the file **directory.txt** only ONCE, and then writes the information to it every time a new file server registers itself with the directory server. Do NOT create 3 individual files!!! Also note that the directory server must first check whether a file by the name of **directory.txt** already exists in the same working directory with the rest of your code or not. If yes, it needs to overwrite the old one with a new one. When your code is tested for a second time, the contents of the **directory.txt** must be overwritten every time, and the new information must NOT be appended to the old information. So, after the process is over the **directory.txt** file could contain the following as an example:

```
File_Server1 22080
File_Server2 23080
File_Server3 24080
```

The directory server expects to hear messages from all the 3 file servers before it receives any messages from the clients. In order to make sure this sequence of events is enforced, we never run the client code before the file server code. However, just to be sure, the directory server should keep track of the file servers that contact it before it can respond to any messages from the clients. The easiest way to avoid such a race condition is to run the directory server first, then the file servers, and after the communication between the file servers and the directory server is over, run the client code. Please do it this way and note that your project will be tested in the same way.

The **directory.txt** must now have 3 lines (2 columns on each line separated by one space) in the same order the file servers contacted the directory server. If you are using `fork( )` and if your code is tested several times, we may see that the file servers are registered in a different order every time, and so every time this info is recorded in the **directory.txt** in a different order.

At the end of this phase, the messages should be displayed on the screen as per the tables provided below.




Figure1: Phase 1, Registration of file servers with directory server

## Phase: 2

In this phase, clients and the directory server communicate using UDP. The static port that should be assigned to each one of them is provided in Table 1. Client 1 should request **doc1** and Client 2 should request **doc2** to the Directory server. Specifically, clients should send a message which includes their Client# followed by file requested. For example, client 1 should send:

Client1 doc1

and client 2 should send:

Client2 doc2

On receipt of these messages the directory server will go through the **resource.txt** to find file servers that have the requested files. If a requested file is present in multiple file servers, then the directory server will lookup the **topology.txt** to determine the file server closest (i.e. the one with less cost) to the client. Directory server will now lookup the **directory.txt** to find the TCP port number associated with the file server and send the client back a message which includes:

File\_Server# TCP\_port\_number

For example if the closest file server was found to be file server 2 and if file server 2 had a TCP port number 23080, then directory server will send the following message to the client:

File\_Server2 23080

This marks the end of phase 2.

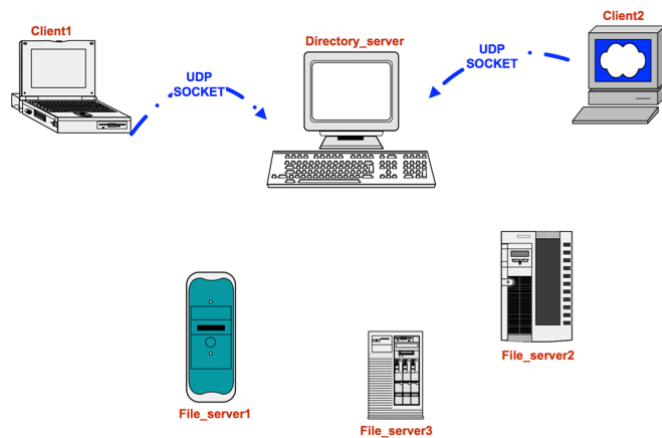


Figure 2: Resource Identification by the Clients

### Phase: 3

In this phase, each file server should open a TCP connection on its respective TCP port number (as mentioned in the Table 1) and should start listening for connections. Then, the clients should establish a TCP connection with the file servers using the TCP port number of the file server that they got from the directory server in phase 2. During this phase, clients will be assigned with a dynamic TCP port number. After the connection is established, the clients send a message which includes the Client# followed by the name of the requested file. For example, client 1 should send:

Client1 doc1

and client 2 should send:

Client2 doc2

Now the file server will send a string **doc#** in response to the client's request. Here **doc#** can either be doc1 or doc2, depending on the file requested by the client. In other words, we assume that there is no actual **download** of the file from the file server, but instead only a string containing the name of the file is sent to the client.

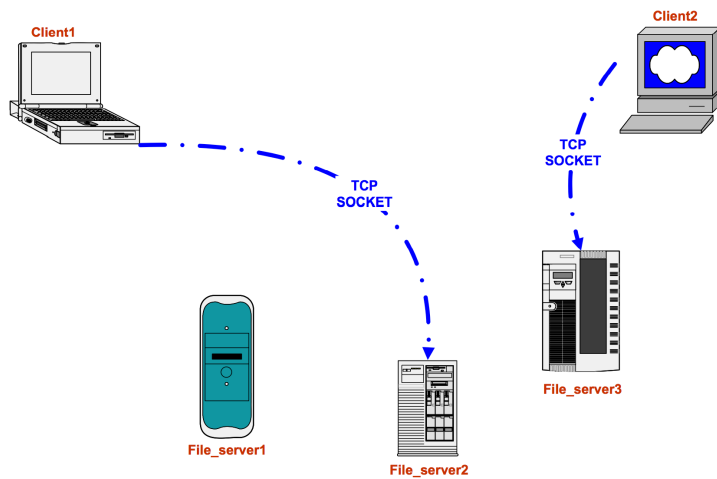


Figure3: File sharing

Table 1. Static and Dynamic assignments for TCP and UDP ports.

Process	Dynamic Ports	Static Ports
Directory Server		1 UDP, 21000 + xxx (last three digits of your ID) (phase 1) 1 UDP, 31000 + xxx (last three digits of your ID) (phase 2)
File Server 1		1 UDP, 22000 + xxx (last digits of your ID) (phase 1) 1 TCP, 41000 + xxx (last digits of your ID) (phase 3)
File Server 2		1 UDP, 23000 + xxx (last digits of your ID) (phase 1) 1 TCP, 42000 + xxx (last digits of your ID) (phase 3)
File Server 3		1 UDP, 24000 + xxx (last digits of your ID) (phase 1) 1 TCP, 43000 + xxx (last digits of your ID) (phase 3)
Client 1	1 TCP (phase 3)	1 UDP, 32000 + xxx (last digits of your ID) (phase 2)
Client 2	1 TCP (phase 3)	1 UDP, 33000 + xxx (last digits of your ID) (phase 2)

## ON-SCREEN MESSAGES

**Table 2. Directory Server on-screen messages**

Event	On-Screen Message
Upon startup of phase 1	Phase 1: The Directory Server has UDP port number ____ and IP address ____.
Upon receiving the registration request information from each of the file server	Phase 1: The Directory Server has received request from File Server #. <i>(you should print this 3 times, one for each file server)</i>
End of phase 1	Phase 1: The directory.txt file has been created. End of Phase 1 for the Directory Server.
Upon startup of phase 2	Phase 2: The Directory Server has UDP port number ____ and IP address ____.
Upon receiving the request from each of the client	Phase 2: The Directory Server has received request from Client #. <i>(you should print this twice, one for each client)</i>
Upon sending the file server details to both the clients	Phase 2: File server details has been sent to Client #. <i>(you should print this twice, one for each client)</i>
End of phase 2	Phase 2: End of Phase 2 for the Directory Server.

**Table 3. File Server 1 on-screen messages**

Event	On-Screen Message
Upon startup of phase 1	Phase 1: File Server 1 has UDP port number ____ and IP address ____.
Upon sending the registration request to the Directory Server	Phase 1: The Registration request from File Server 1 has been sent to the Directory Server.
End of Phase 1	Phase 1: End of Phase 1 for File Server 1.
Upon startup of phase 3	Phase 3: File Server 1 has TCP port ____ and IP address ____.
Upon receiving the request from	Phase 3: File Server 1 received the request from the



the client	<clientname> for the file <filename>. <i>(you should print this as many times as the requests received by the file server)</i>
Upon transferring the file to the client (i.e. sending the string containing the requested filename)	Phase 3: File Server 1 has sent <filename> to <clientname>. <i>(you should print this as many times as the requests received by the file server)</i>

**Table 4. File Server 2 on-screen messages**

Event	On-Screen Message
Upon startup of phase 1	Phase 1: File Server 2 has UDP port number ____ and IP address ____.
Upon sending the registration request to the Directory Server	Phase 1: The Registration request from File Server 2 has been sent to the Directory Server.
End of Phase 1	Phase 1: End of Phase 1 for File Server 2.
Upon startup of phase 3	Phase 3: File Server 2 has TCP port ____ and IP address ____.
Upon receiving the request from the client	Phase 3: File Server 2 received the request from the <clientname> for the file <filename>. <i>(you should print this as many times as the requests received by the file server)</i>
Upon transferring the file to the client (i.e. sending the string containing the requested filename)	Phase 3: File Server 2 has sent <filename> to <clientname>. <i>(you should print this as many times as the requests received by the file server)</i>

A A 10/28/13 12:25 PM

Deleted: End of phase 3

... [1]

**Table 5. File Server 3 on-screen messages**

Event	On-Screen Message
Upon startup of phase 1	Phase 1: File Server 3 has UDP port number ____ and IP address ____.
Upon sending the registration request to the Directory Server	Phase 1: The Registration request from File Server 3 has been sent to the Directory Server.
End of Phase 1	Phase 1: End of Phase 1 for File Server 3.

A A 10/28/13 12:25 PM

Deleted: End of phase 3

... [2]

Upon startup of phase 3	Phase 3: File Server 3 has TCP port _____ and IP address _____.
Upon receiving the request from the client	Phase 3: File Server 3 received the request from the <clientname> for the file <filename>. <i>(you should print this as many times as the requests received by the file server)</i>
Upon transferring the file to the client (i.e. sending the string containing the requested filename)	Phase 3: File Server 3 has sent <filename> to <clientname>. <i>(you should print this as many times as the requests received by the file server)</i>

**NOTE:**

<filename> : file requested by the client. Can be either doc1 or doc2.

<clientname> : Name of the client sending the request . Can be either Client 1 or Client 2.

A A 10/28/13 12:25 PM

Deleted: End of phase 3

... [3]

**Table 6. Client 1 on-screen messages**

Event	On-Screen Message
Upon startup of phase 2	Phase 2: Client 1 has UDP port number _____ and IP address _____.
When the file request is sent to the directory server	Phase 2: The File request from Client 1 has been sent to the Directory Server.
When the response from directory server is received	Phase 2: The File <u>requested by Client 1</u> is present in <file_server_name> and the File Server's TCP port number is <received_tcp_port_number>.
End of phase 2	Phase 2: End of <u>Phase 2</u> for Client 1.
Upon startup of phase 3	Phase 3: Client 1 has dynamic TCP port number _____ and IP address _____.
When the file request is sent to the file server	Phase 3: The File request from Client 1 has been sent to the <file_server_name>
When the file is received by client 1	Phase 3: <u>Client 1</u> received <u>&lt;filename&gt;</u> from <file_server_name>.
End of phase 3	Phase 3: End of Phase 3 for Client 1.

**Table 7. Client 2 on-screen messages**

Event	On-Screen Message
Upon startup of phase 2	Phase 2: Client 2 has UDP port number ____ and IP address ____.
When the file request is sent to the directory server	Phase 2: The File request from Client 2 has been sent to the Directory Server.
When the response from directory server is received	Phase 2: The File <u>requested by Client 2</u> is present in <file_server_name> and the File Server's TCP port number is <received_tcp_port_number>.
End of phase 2	Phase 2: End of <u>Phase 2</u> for Client 2.
Upon startup of phase 3	Phase 3: Client 2 has dynamic TCP port number ____ and IP address ____.
When the file request is sent to the file server	Phase 3: The File request from Client 2 has been sent to the <file_server_name>
When the file is received by client 2	Phase 3: <u>Client 2</u> received <u>&lt;filename&gt;</u> from <file_server_name>.
End of phase 3	Phase 3: End of Phase 3 for Client 2.

**NOTE:**

<file\_server\_name>: file server returned back to the client by directory server. Can be File Server 1, File Server 2, or File Server 3.

<file\_server\_tcp\_port\_number>: TCP port number of the file server returned by the directory server.

<filename> : file requested by the client. Can be either doc1 or doc2.

|

